

# Linux Kernel Data Structure

---

Chung-Ang University



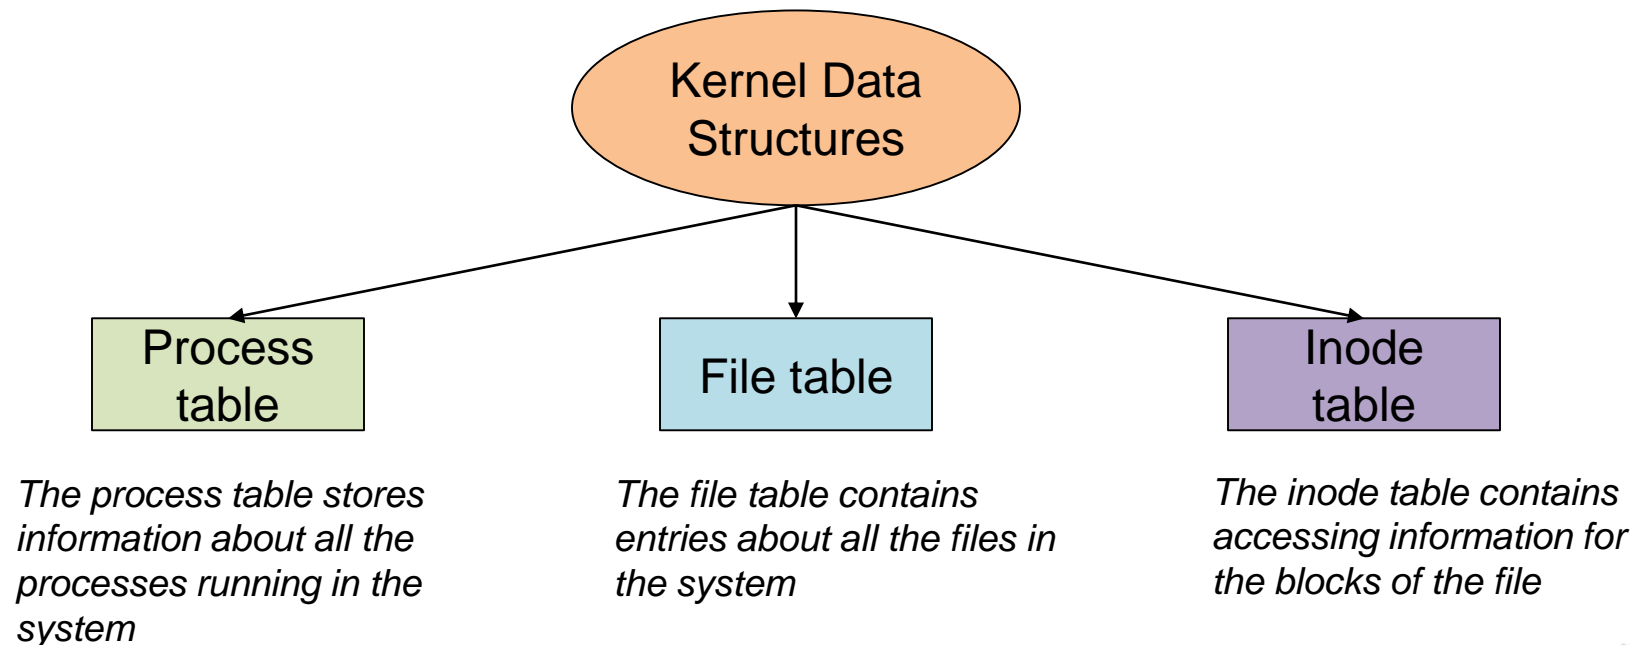
# Kernel Data Structure

- The kernel data structures are very important as they store data about the current state of the system
  - For example, if a new process is created in the system, a kernel data structure is created that contains the detail information or state about the process
  - Kernel data structures are only accessible by the kernel and its subsystems
  - They contain data as well as pointers to other data structures



# Kernel Data Structure

- The kernel data structure stores and organizes a lot of information
  - For example, it has data about which processes are running in the system, their memory requirements, files in use, etc
  - To handle all this, three important structures are used
  - These are process table, file table, and inode table



# Kernel Data Structures

- This chapter introduces several built-in data structures for use in Linux kernel code
  - As with any large software project, the Linux kernel provides these generic data structures to encourage code reuse
  - Kernel developers can use these data structures



# Types of Linux Kernel Data structure


- Linked list
- Red Black tree
- Hash table



# Singly linked list

- Singly linked list is a basic linked list type
- Singly linked list is a collection of nodes linked together in a sequential way
- Each node of singly linked list contains a data field and an address field which contains the reference of the next node

```
struct my_list_element{  
    void *data; /* void pointer to point on a generic data */  
    struct my_list_element *next; /* pointer to a next element */  
};
```




- Starts from **HEAD** and terminates at **NULL**
- Traverses forward only
- When empty, **HEAD** is **NULL**



# Doubly linked list

- Each node contains three fields: two link fields (references to the previous and to the next node in the sequence of nodes) and one data field




```
struct my_list_element{  
    void *data; /* void pointer to point on a generic data */  
    struct my_list_element *prev; /* pointer to a previous element */  
    struct my_list_element *next; /* pointer to a next element */  
};
```

- Starts from **HEAD** and terminates at **NULL**
- Traverses forward and backward
- When empty, **HEAD** is **NULL**




# Circular linked list

- Circular linked list is a linked list where all nodes are connected to form a circle
  - There is no NULL at the end
  - A circular linked list can be a singly circular linked list or doubly circular linked list



**Singly circular linked list**



**Doubly circular linked list**


- Starts from **HEAD** and terminates at **HEAD**
- When empty, **HEAD** is **NULL**





# Linux Kernel Linked List

- Linux kernel linked list

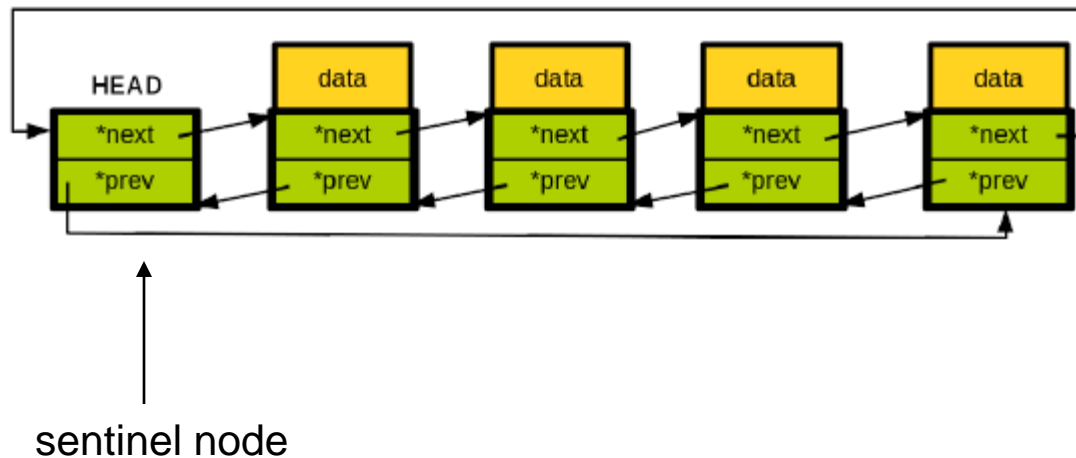


- Starts from HEAD and terminates at HEAD
- When empty, HEAD is not NULL
  - Prev and next of HEAD points HEAD
  - HEAD is a sentinel node
- Easy to insert a new element at the end of a list
- There is no exceptional case to handle **NULL**



# Linux Kernel Linked List

- Linux kernel linked list is a circular doubly linked list
- Two differences from the typical design
  - Embedding a linked list node in the structure
  - Using a sentinel node as a list header
- `linux/include/linux/list.h`



# Linux Kernel Linked List

- struct list\_head is the key data structure
- list\_head is embedded in the data structure
- Start of a list is also list\_head my\_car\_list → sentinel node

```
struct list_head { /* kernel linked list data structure */
    struct list_head *next, *prev;
};

struct car {
    struct list_head list; /* add list_head instead of prev and next */
    unsigned int max_speed; /* put data directly */
    unsigned int drive_when_num;
    unsigned int price_in_dollars;
};

struct list_head my_car_list; /* head is also list_head */
```



# Linux Kernel Linked List

- Getting a data from **list\_head**
  - use `list_entry(ptr, type, member)`

```
struct car *amazing_car = list_entry(car_list_ptr, struct car, list)
```

`list_entry` – get the struct for this entry

@ptr: the &struct `list_head` pointer

@type: the type of the struct this is embedded in

@member: the name of the `list_head` within the struct

#define **list\_entry**(ptr, type, member) **container\_of**(ptr, type, member)



# Linux Kernel Linked List

## ■ Linux kernel linked list APIs

```
/* Insert a new entry after the specified head */
void list_add(struct list_head *new, struct list_head *head);

/* Insert a new entry before the specified head */
void list_add_tail(struct list_head *new, struct list_head *head);

/* Delete a list entry
 * NOTE: You still have to take care of the memory deallocation if
 * needed */
void list_del(struct list_head *entry);

/* Delete from one list and add as another's head */
void list_move(struct list_head *list, struct list_head *head);

/* Delete from one list and add as another's tail */
void list_move_tail(struct list_head *list, struct list_head *head);

/* Join two lists (merge a list to the specified head) */
void list_splice(const struct list_head *list, struct list_head *head);
```



# Linux Kernel Linked List

## ■ Linux kernel linked list APIs

```
/**
 * list_for_each - iterate over a list
 * @pos: the &struct list_head to use as a loop cursor.
 * @head: the head for your list.
 */
#define list_for_each(pos, head) \
    for (pos = (head)->next; pos != (head); pos = pos->next)

/**
 * list_for_each_entry - iterate over list of given type
 * @pos: the type * to use as a loop cursor.
 * @head: the head for your list.
 * @member: the name of the list_head within the struct.
 */
#define list_for_each_entry(pos, head, member) \
    for (pos = list_first_entry(head, typeof(*pos), member); \
         &pos->member != (head); \
         pos = list_next_entry(pos, member))
```



# Linux Kernel Linked List

- An example of Linux kernel linked list

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/init.h>
#include <linux/list.h>
#include <linux/slab.h> // for kmalloc

struct my_node {
    struct list_head list;
    int data;
};
```

```
int __init hello_module_init(void)
{
    struct_exmaple();
    printk("module init\n");
    return 0;
}

void __exit hello_module_cleanup(void)
{
    printk("Bye Module\n");
}

module_init(hello_module_init);
module_exit(hello_module_cleanup);
```



# Linux linked list

## ■ Linux linked list example

```
void struct_exmample(void)
{
    struct list_head my_list;

    /* initialize list */
    INIT_LIST_HEAD(&my_list);

    /* add list element */
    int i;
    for (i = 0; i < 10; i++) {
        struct my_node *new = kmalloc(sizeof(struct my_node), GFP_KERNEL);
        new->data = i;
        list_add(&new->list, &my_list);
    }

    struct my_node *current_node; /* This will point on the actual data structures during the iteration */
    struct list_head *p; /* Temporary variable needed to iterate; */
    list_for_each(p, &my_list) {
        current_node = list_entry(p, struct my_node, list);
        printk("current value : %d\n", current_node->data);
    }
    list_for_each_entry(current_node, &my_list, list) {
        printk("current value : %d\n", current_node->data);
    }
}
```

```
struct my_node {
    struct list_head list;
    int data;
};
```

```
#define list_for_each(pos, head) \
    for (pos = (head)->next; pos != (head); pos = pos->next)

#define list_for_each_entry(pos, head, member) \
    for (pos = list_first_entry(head, typeof(*pos), member); \
         &pos->member != (head); \
         pos = list_next_entry(pos, member))
```





# Linux linked list

```
/* iterate over a list reversely*/
list_for_each_entry_reverse(current_node, &my_list, list){
    printk("current value : %d\n", current_node->data);
}

/* delete list element */
list_for_each_entry(current_node, &my_list, list){
    if(current_node->data == 2){
        printk("current node value : %d\n", current_node->data);
        list_del(&current_node->list);
        kfree(current_node);
    }
}

/* iterate over a list */
list_for_each_entry(current_node, &my_list, list){
    printk("current value : %d\n", current_node->data);
}
}
```




# Linux Kernel Linked List

- **Usage of linked list in the kernel**
  - Kernel code makes extensive use of linked lists
  - a list of threads under the same parent PID
  - a list of superblocks of a file system
  - and many more




# Red-Black Tree

- Tree basics : binary tree
- Properties
  - Nodes have zero, one, or two children
  - Root has no parent, other nodes have one




# Red-Black Tree

- Tree basics : binary search tree
- Properties
  - Left children  $<$  parent
  - Right children  $>$  parent
  - Search and ordered traversal are efficient




# Red-Black Tree

- Tree basics : balanced binary search tree
- Properties
  - Depth of all leaves differs by at most one



# Red-Black Tree

- Tree basics : red-black tree
- Self-balancing binary search tree
- Properties
  - A type of self-balancing binary search tree
  - Nodes: red or black
  - Leaves: black, no data



# Linux Kernel Red-Black Tree

- Self-balancing binary search tree


/\* Rbtree node, which is embedded to your data structure like

\* list\_head and hlist node \*/

```
struct rb_node {  
    unsigned long __rb_parent_color;  
    struct rb_node *rb_right;  
    struct rb_node *rb_left;  
};
```

/\* Root of a rbtree \*/

```
struct rb_root {  
    struct rb_node *rb_node;  
};
```



# Linux Kernel Red-Black Tree

## ■ red-black tree APIs

```
/* Find logical next and previous nodes in a tree */
```

```
struct rb_node *rb_next(const struct rb_node *);
```

```
struct rb_node *rb_prev(const struct rb_node *);
```

```
struct rb_node *rb_first(const struct rb_root *);
```

```
struct rb_node *rb_last(const struct rb_root *);
```

```
/* Insert a new node under a parent connected via rb_link */
```

```
void rb_link_node(struct rb_node *node, struct rb_node *parent,  
                 struct rb_node **rb_link);
```

```
/* Re-balance an rbtree after inserting a node if necessary */
```

```
void rb_insert_color(struct rb_node *, struct rb_root *);
```

```
/* Delete a node */
```

```
void rb_erase(struct rb_node *, struct rb_root *);
```





# Linux Kernel Red-Black Tree

## ■ Red-Black Tree example

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/init.h>
#include <linux/rbtree.h>
#include <linux/slab.h> // for kmalloc
```

```
#define FALSE 0
#define TRUE 1
```

```
struct my_node {
    struct rb_node node;
    int key;
    int value;
};
```

```
int __init hello_module_init(void)
{
```

```
    RB_exmaple();
    printk("module init\n");
    return 0;
```

```
}
```

```
void __exit hello_module_cleanup(void)
{
```

```
    printk("Bye Module\n");
```

```
}
```

```
module_init(hello_module_init);
module_exit(hello_module_cleanup);
```



# Linux Kernel Red-Black Tree

```
void RB_exmaple(void)
{
    struct rb_root root_node = RB_ROOT;
    int i;

    /* rb_node create and insert */
    for(i=0;i<20;i++){
        struct my_node *new_node = kmalloc(sizeof(struct my_type),GFP_KERNEL);
        if(!new_node)
            return NULL;
        new_node->value = i*10;
        new_node->key = i;
        ret = rb_insert_color(new_node, &root_node);
    }

    /* rb_tree traversal using iterator */
    struct rb_node *iter_node;
    for (iter_node = rb_first(&root_node); iter_node; iter_node = rb_next(iter_node))
        printk("(key,value)=(%d.%d)\n", \
            rb_entry(iter_node, struct my_node, node)->key, \
            rb_entry(iter_node, struct my_node, node)->value);

    /* rb_tree delete node */
    rb_erase(new_node, &root_node);
}
```



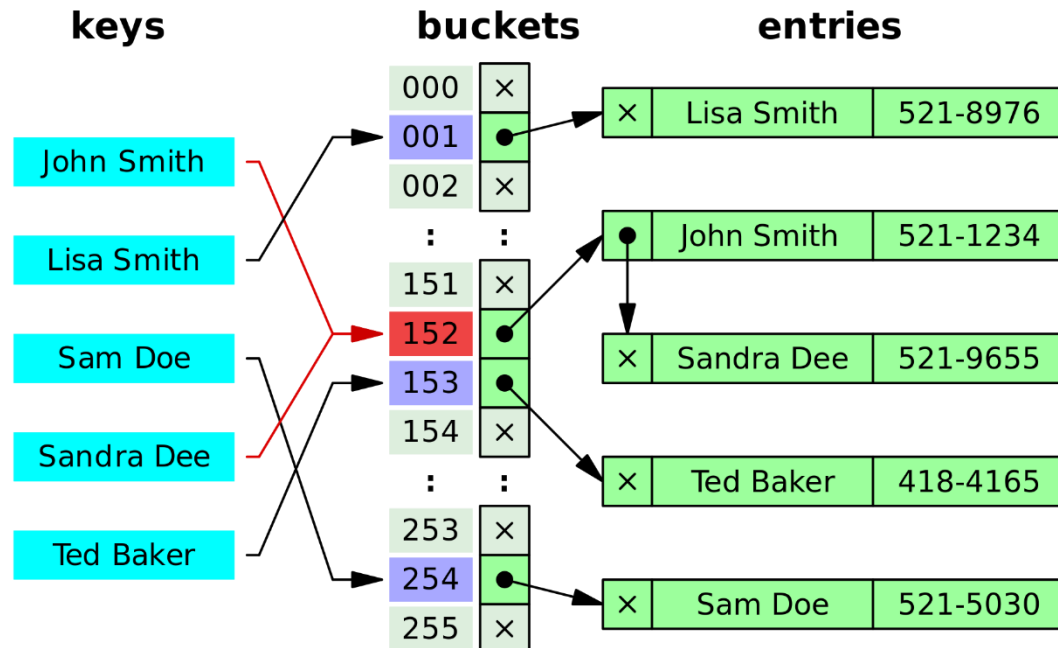
# Red-Black Tree

- **Usage of Red-Black tree in the kernel**
  - Completely Fair Scheduling (CFS)
    - Default task scheduler in Linux
    - Each task has **vruntime** , which presents how much time a task has run
    - CFS always picks a process with the smallest **vruntime** for fairness
    - Per-task **vruntime** structure is maintained in a rbtree



# Hash table

- A simple fixed-size chained hash table
  - The size of bucket array is fixed at initialization as a  $2^N$
  - Each bucket has a singly linked list or doubly linked list to resolve hash collision



# Hash table

## ■ Hash table APIs

```
/**
```

```
 * Define a hashtable with 2^bits buckets
```

```
 */
```

```
#define DEFINE_HASHTABLE(name, bits) ...
```

```
/**
```

```
 * hash_init - initialize a hash table
```

```
 * @hashtable: hashtable to be initialized
```

```
 */
```

```
#define hash_init(hashtable) ...
```

```
/**
```

```
 * hash_add - add an object to a hashtable
```

```
 * @hashtable: hashtable to add to
```

```
 * @node: the &struct hlist_node of the object to be added
```

```
 * @key: the key of the object to be added
```

```
 */
```

```
#define hash_add(hashtable, node, key) ...
```



# Hash table

## ■ Hash table APIs

```
/**
 * hash_for_each - iterate over a hashtable
 * @name: hashtable to iterate
 * @bkt: integer to use as bucket loop cursor
 * @obj: the type * to use as a loop cursor for each entry
 * @member: the name of the hlist_node within the struct
 */
#define hash_for_each(name, bkt, obj, member) ...

/**
 * hash_del - remove an object from a hashtable
 * @node: &struct hlist_node of the object to remove
 */
void hash_del(struct hlist_node *node);
```



# Hash table

## ■ Hash table example

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/init.h>
#include <linux/hashtable.h>
#include <linux/slab.h> // for kcalloc

#define MY_HASH_BITS 2

struct my_node {
    u32 key;
    int value;
    struct hlist_node hnode;
};
```

```
int __init hello_module_init(void)
{
    hash_exmample();
    printk("module init\n");
    return 0;
}

void __exit hello_module_cleanup(void)
{
    printk("Bye Module\n");
}

module_init(hello_module_init);
module_exit(hello_module_cleanup);
```



# Hash table

```
void hash_exmaple(void)
{
    DEFINE_HASHTABLE(my_hash, MY_HASH_BITS);
    hash_init(my_hash);

    /* (key,value) insert */
    int i;
    for(i=0;i<10;i++){
        struct my_node *new = kmalloc(sizeof(struct my_node), GFP_KERNEL);
        new->value = i * 10;
        new->key = i;
        memset(&new->hnode, 0, sizeof(struct hlist_node));
        hash_add(my_hash,&new->hnode,new->key);
    }

    /* You can complete hash_for_each() function!!! */
    hash_for_each(.....){
        .....
    }

    /* You can complete hash_del() function!!! */
    hash_del(.....){
        .....
    }
}
```





# Hash table

- Usage of hash table in the kernel
  - Hugepage
    - finds physically consecutive 4KB pages
    - remaps consecutive 4KB pages to a 2MB page (huge page)
    - saves TLB entries and improves memory access performance by reducing TLB miss
    - maintains per-process memory structure, **struct mm\_struct**

